

# Executive Summary

This document outlines the deployment plan for an AI-powered lead enrichment and scoring pipeline across EliseAI's Sales Development team. The system replaces 45 manual research per lead with a fully automated enrichment run improving both coverage and outreach quality.

The rollout follows a five-phase model: internal validation, pilot, A/B controlled expansion, full deployment, and continuous improvement. Each phase has explicit go/no-go gates to protect SDR experience and data quality throughout.

# Rollout Timeline

Eight weeks from first test to full-team adoption, with a continuous improvement track from month three onwards.

|                                 |          |
|---------------------------------|----------|
| Phase 1: Pre-Launch Validation  | Week 1-2 |
| Phase 2: Pilot                  | Week 2-3 |
| Phase 3: A/B Shadow Rollout     | Week 3-4 |
| Phase 4: Full Deployment        | Week 4-5 |
| Phase 5: Continuous Improvement | Month 1+ |

---

Phase 1 · Week 1

## Pre-Launch Validation

---

Before rolling anything out to the team, we validate against real data. If the scoring logic doesn't hold up against historical conversion outcomes, we fix it here not in production.

### Historical Data Backfill

- Pull last 100 inbound leads from CRM (past 30 days)
- Run through enrichment pipeline and compare AI tier assignments to actual conversion outcomes

**Success metric:** ≥70% of Tier A leads converted to opportunity

### API Reliability Testing

- Stress test: 100 leads processed simultaneously

- Monitor API success rates across all data sources (target: ≥85% per source)
- Test failure handling: bad addresses, missing fields, partial data

**Success metric:** Zero pipeline crashes; graceful degradation on any API failure

### Output Quality Audit

- Manual review of 20 AI-generated outreach emails
- Each email must reference ≥2 specific enriched data points with actionable talking points

**Success metric:** ≥90% of emails rated send-ready (zero or light SDR edits needed)

**Deliverable:** System validation report accuracy, speed, and quality benchmarks documented

---

Phase 2 · Week 2-3

## Pilot with lead/high performers

---

Small enough to debug fast. Large enough to validate across different lead types. High performers give quality signal, not just complaints.

### SDR Selection

- 2 inbound specialists (web form leads) + 2 outbound specialists (conference / referral leads)
- Tests the system across different lead sources, data quality levels, and working styles

### Integration Setup

- New lead hits CRM → Zapier webhook fires → enrichment runs automatically
- Results populate custom CRM fields: Tier, Score, Email Draft, Key Signals
- Tier A Slack DM fires within minutes of lead arrival

### What We Measure — Daily for 2 Weeks

- Time saved per lead vs control group benchmark
- Email reply rate: AI-drafted vs manually written
- SDR satisfaction: daily 1–10 score
- Tier accuracy: did Tier A leads actually book meetings?

### Go Criteria

- Time savings ≥20 minutes per lead

- Reply rate maintained or improved vs manual
- SDR satisfaction  $\geq 7/10$
- Zero critical bugs

**Deliverable:** Pilot results deck with metrics, SDR testimonials, and refinement log

---

Phase 3 · Week 3-4

## Shadow A/B Rollout

---

A control group lets us prove the tool is driving results not just automation. It also de-risks scaling by surfacing infrastructure issues before 100% of the team is affected.

### Team Structure

- Group A (Tool Users): 50% of SDRs, randomly assigned
- Group B (Control): 50% continue manual workflow
- Stratified by tenure, geography, and lead volume for balanced comparison

### Expanded Infrastructure

- Scheduler runs every 30 minutes new leads processed automatically
- Live dashboard for SDR Manager: tier distribution, API health, throughput
- Weekly digest: leads processed, Tier A count, avg time saved per lead

### Primary Comparison Metrics

- Leads contacted within 24 hours (%)
- Email reply rate (%)
- Meetings booked per week
- Research vs conversations time split — tracked via CRM activity logs

### Go Criteria

- Group A contacts  $\geq 20\%$  more leads vs Group B
- Group A reply rate  $\geq$  Group B
- No bugs affecting more than 5% of processed leads

**Deliverable:** A/B results report: quantified business impact across time, quality, and pipeline

---

Phase 4 · Week 4-5

## Full Deployment

---

### Onboarding

- 45-minute live training session recorded for async viewing
- "How to use AI SDR outputs" playbook linked from CRM
- Daily 15-minute office hours during the first week

### Production Infrastructure

- Migrate from Zapier webhook to native CRM integration (Salesforce / HubSpot)
- Scale backend: additional Railway dynos + Redis caching for API responses
- Monitoring: Datadog / Sentry for uptime and error tracking
- On-call rotation for first two weeks: GTM Engineer + 1 backend engineer

### Long-Term Success Metrics

- Adoption rate: >90% of SDRs actively using tool (≥5 leads / week)
- Coverage: >95% of inbound leads enriched within 1 hour
- Email acceptance: >80% used as-is or with light edits
- Business impact: increase in qualified meetings booked per SDR per week

**Deliverable:** Production launch checklist complete, full team trained, monitoring dashboards live

---

Phase 5 · Month 1+

## Continuous Improvement

---

### Adjacent Use Cases

- AE Handoff Briefs: auto-generated account research for closing reps
- CS Expansion Triggers: flag growth signals funding, headcount, leadership changes
- Event Follow-Up: enrich trade show and conference leads in real time

### Data Feedback Loop

- Track which enriched signals correlate with deal close

- Retrain routing logic monthly remove APIs with low predictive value
- Feed closed-won data back into scoring model quarterly

### Governance

- Monthly: GTM Engineer + SDR Manager + RevOps review
- Quarterly: Sales Leadership approves new features and roadmap
- Continuous: SDRs submit feature requests via Slack poll

## Decision Gates

Each gate is binary. Criteria not met means stop, fix, and re-evaluate. No gate-skipping.

### Gate 1: After Pilot (End of Week 3)

| GO                                | NO-GO                     |
|-----------------------------------|---------------------------|
| Time savings $\geq 20$ min/lead   | Critical bugs unresolved  |
| SDR satisfaction $\geq 7/10$      | SDR satisfaction $< 5/10$ |
| Reply rate maintained or improved | Reply rate drops $> 10\%$ |

### Gate 2: After A/B Test (End of Week 6)

| GO                                              | NO-GO                                |
|-------------------------------------------------|--------------------------------------|
| Group A outperforms Group B on $\geq 2$ metrics | No measurable improvement vs control |
| No scaling issues observed                      | API costs exceed approved budget     |
| API costs within budget                         | Scaling issues unresolved            |

## Risk Mitigation

| Risk                              | Impact | Mitigation                                                      |
|-----------------------------------|--------|-----------------------------------------------------------------|
| API rate limits during peak hours | High   | Request queuing, backup API keys, 24-hour response caching      |
| LLM generates inappropriate email | High   | Human-in-loop for first 500 emails, automated compliance filter |
| SDRs don't trust AI scores        | Medium | Show per-tier reasoning; allow SDR override with logging        |

| Risk                             | Impact | Mitigation                                                     |
|----------------------------------|--------|----------------------------------------------------------------|
| CRM integration breaks           | High   | Fallback to Google Sheets, 24h SLA on fixes, monitoring alerts |
| High lead volume causes slowdown | Medium | Auto-scaling backend, async processing queue, load benchmarks  |
| Data privacy / security concerns | High   | SOC 2 review, encrypted API keys, audit logs, retention policy |

## Stakeholder Map

| Stakeholder      | Role              | Cadence                    | Key Concerns                     |
|------------------|-------------------|----------------------------|----------------------------------|
| SDR Manager      | Primary owner     | Weekly 1:1s                | Team productivity, adoption rate |
| RevOps Lead      | Data & systems    | Bi-weekly sync             | CRM integration, data quality    |
| IT / Security    | Compliance        | As-needed                  | API security, data handling      |
| Sales Leadership | Executive sponsor | Monthly review             | ROI, pipeline impact             |
| Individual SDRs  | End users         | Daily Slack + office hours | Ease of use, output quality      |
| GTM Engineer     | Builder & owner   | Daily monitoring           | Reliability, iteration velocity  |

## Budget & Resources

### One-Time Costs

- API subscriptions: \$0 (free tiers sufficient for MVP)
- Deployment infrastructure (Railway, Netlify): \$0–\$50/month
- Training materials: ~8 hours GTM Engineer time

### Ongoing Monthly Costs at Scale

- API usage (Census, WalkScore, FRED, NewsAPI): ~\$100–\$200/month
- LLM inference (Groq): \$0 on free tier; \$50–\$100/month at full scale
- Backend hosting: \$20–\$50/month
- Monitoring tools: \$0–\$30/month

### Team Time Investment

- GTM Engineer: 40 hrs/week for month 1; 10 hrs/week ongoing
- SDR Manager: 5 hrs/week during rollout
- RevOps: 2 hrs/week for CRM integration support

## **Conclusion:**

SDRs should be having conversations, not doing research. This pipeline handles the research every lead, every time, in under 90 seconds.